

Evidence-Backed KPI Decision Engine & Causal Governance Architecture

python

3.11

FastAPI

0.111.0

React

19.2.8

Vite

8.2.2

PostgreSQL

15

ChromaDB

0.5.0

License

Proprietary

1. Executive Overview

The Problem

- Modern enterprise observability and Business Intelligence dashboards suffer from a fundamental epistemic disconnect:
- Dashboards Show *What* Changed, Not *Why*:** Traditional BI tools display metric drops (e.g. conversion down 14%, latency up 300%) but cannot isolate whether the root cause was an internal deployment, an upstream payment gateway outage, a competitor marketing campaign, or an ETL ingestion telemetry artifact.
 - Generic LLMs Hallucinate Causal Chains:** Standard AI assistants generate convincing qualitative stories but lack deterministic mathematical grounding. When presented with correlation, they jump to conclusions, propose ungrounded production interventions (e.g. rolling back software during a third-party ISP outage), and ignore enterprise governance boundaries.
 - Absence of Epistemic Guards & Decision Rights:** Existing systems lack mathematical abstention mechanisms. When data is sparse (<14 days), ambiguous, or corrupted, generic AI guesses rather than abstaining.

The Solution


BusinessIntelligence.ai is an evidence-backed KPI decision engine designed for enterprise reliability engineering, financial operations, and executive decision-making.

The platform separates **deterministic mathematical auditing and content-hashed evidence retrieval** from **LLM qualitative hypothesis synthesis and executive briefing**. Every candidate hypothesis is put on trial against 5 deterministic non-LLM verification rules and hard epistemic gates. The system enforces strict role-based data entitlements (`Analyst` , `CFO` , `Manager`), adheres to governed operational levers, and triggers automated safety abstentions whenever evidence confidence is marginal or ambiguous.

2. Solution Architecture: The 9-Engine Investigation Pipeline

The platform processes incidents through a strictly governed 9-stage pipeline ($E_1 \rightarrow E_9$):

9-ENGINE INVESTIGATION PIPELINE



Why This Architecture?

The architecture is built on a core principle: **never ask an LLM to perform mathematical auditing, and never force a database to formulate qualitative domain context.**

Quantitative truth	→ Relational SQL / Rolling Statistics / Corridor Math
Qualitative explanation	→ Large Language Models (LLM) under strict schemas
Causal verification	→ Deterministic 5-rule audit & weakest-link scoring
Operational decision	→ Policy rules + Role entitlements + LLM briefing
Expected outcome	→ Calibrated parametric simulation
Institutional learning	→ Vector retrieval + Human feedback lifecycle

Detailed Engine Breakdown

Engine	Name	Primary Responsibility	Input	Output	Processing Method	Governance & Safety Safeguards
E1	KPI Store	Ingests time-series telemetry and computes historical rolling baseline corridors (μ, σ).	KPI Contract YAML, Time Window [t_0, t_1], Scope filter	Normalized Time-Series DataFrames	Deterministic SQL / Pandas	Statistical Cold-Start Guard (flags datasets with < 14 intervals of baseline history).
E2	Signal Engine	Detects statistical anomalies (z-score) and computes revenue materiality rankings. E1 Time-Series Data Prioritized `list[AnomalySignal]` with materiality tiers **Deterministic Statistical Math** Adverse direction filtering; data quality anomaly detection (ETL ingestion delay vs real drop). **E3** **Diagnostic Engine** Decomposes material anomalies across dimensional slices (`device`, `region`, `channel`). E2 Top Prioritized Anomaly, SQL Transaction Store `list[DimensionContribution]` (cohort contribution %) **Deterministic SQL GROUP BY Aggregation** Multi-dimensional variance isolation; zero-segment safe handling without crashing. **E4** **Evidence Dossier** Assembles verifiable telemetry records, deployment logs, and vendor feeds. E2/E3 Signals, Persona Entitlements YAML `list[Evidence]` with content hashes and timestamps **Hybrid Vector (ChromaDB) + Relational Entitlement Filter** Strict role-based entitlement filtering (masks SRE/deployment logs from non-technical personas). **E5** **Hypothesis Studio** Formulates candidate causal				

Engine	Name	Primary Responsibility	Input	Output	Processing Method	Governance & Safety Safeguards
		explanations structured under a 4-layer ontology. E2 Signals, E3 Contributions, E4 Evidence Dossier 'list[Hypothesis]' with mechanisms & evidence citations **LLM Synthesis (Groq Qwen / Ollama Llama)** Zero-number qualitative propositions; prompt-enforced prohibition of ungrounded citation hashes. **E6** **Challenge Engine** Audits candidate hypotheses against 5 deterministic rules and epistemic gates. E5 Hypotheses, E4 Evidence Records, Anomaly Windows 'list[ScoredHypothesis]' with 'AuditVerdict' & rule scores **100% Deterministic Pure Mathematical Functions** **Root-Cause Evidence Gate:** Mandatory discriminative evidence for release/provider claims; Weakest-link constraint. **E7** **Decision Engine** Synthesizes governed operational action directives or triggers safety guards. E6 Audited Scorecards, Entitlements, Controllable Levers 'DecisionPayload' with action directive & persona narrative **Deterministic Policy Rules + LLM Executive Briefing** **Fail-Closed Abstention:** Automated mitigation suppressed if score ≤ 0.70, margin < 0.15, or lever unauthorized.				
E8	Outcome Simulator	Projects recovery trajectories and Mean Time to Normalcy (MTTN) under proposed actions.	E7 Action Directive, E2 Anomaly Baseline Delta	OutcomeProjection with recovery curves $y(t)$	Deterministic Calibrated Exponential Decay Simulation	Explicit [SIMULATED] provenance tag; non-remedial actions suppress recovery projections; causal disclaimer.

Engine	Name	Primary Responsibility	Input	Output	Processing Method	Governance & Safety Safeguards
E9	Memory Engine	Embeds verified incident cases and retrieves lifecycle-filtered institutional precedents.	E7 Decision, Analyst Feedback Record, ChromaDB	Historical precedent citations & matching score	Vector Embeddings (ChromaDB / Cosine Similarity)	5-State Precedent Lifecycle: Validated precedents receive ranking boosts; disputed/suppressed records are excluded.

3. LLM vs. Non-LLM Processing Matrix

To guarantee mathematical reproducibility and enterprise auditability, the system strictly isolates stochastic LLM generation from deterministic computation:

Pipeline Stage	Method	LLM Required?	Exact Role of LLM vs. Deterministic Code
E1 KPI Store	Relational SQL & Rolling Math	✗ No	100% deterministic database queries and statistical baseline computation (μ, σ) .
E2 Signal Detection	Statistical z\$-Score & Materiality ✗ **No** 100% deterministic formula: $z = (x_t - \mu) / \sigma$; Materiality = $W_{rev} \times \Delta$.		
E3 Diagnostic Decomposition	Relational Dimensional Slicing	✗ No	100% deterministic SQL group-by aggregation across device, region, and channel cohorts.
E4 Evidence Dossier	Hybrid Vector + Relational Security	⚠ Embeddings Only	Dense embedding inference for semantic search; 100% deterministic entitlement masking.
E5 Hypothesis Studio	Causal Ontology Formulation	✅ Yes	LLM generates qualitative explanations conforming to the 4-layer causal schema.
E6 Challenge Audit	Deterministic Constraint Scoring	✗ No	Zero LLM involvement in audit scoring. Evaluated via 5 non-LLM Python rules and hard gates.
E7 Decision Engine	Governed Action Formulation	⚠ Hybrid	Deterministic rule decides VERIFIED vs ABSTAIN ; LLM synthesizes role-adapted narrative.
E8 Outcome Simulation	Exponential Recovery Decay	✗ No	100% deterministic parametric simulation: $y(t) = y_{target} + (y_0 - y_{target})e^{-\lambda t}$.
E9 Memory Engine	Precedent Vector Store	⚠ Embeddings Only	Dense embedding generation; deterministic cosine/HNSW matching and 5-state lifecycle filtering.

Early Deterministic Guards (LLM Bypass)

To optimize LLM economics and eliminate hallucinations, deterministic early guards can bypass LLM inference entirely for scenarios that fail data-quality, baseline, or nominal-corridor checks:

* **Cold-Start Guard (< 14 Days):** If baseline history is insufficient, hypothesis generation is skipped (0\$ LLM calls). * **Nominal Corridor Guard ($|z| < 3.0\sigma$):** * *
If all telemetry streams fluctuate normally within calibrated corridor bounds, the system declares 'SYSTEM NOMINAL' (0\$ LLM calls).
* **Data Quality Guard (ETL Delay):** If revenue drops without application error or gateway latency elevation, the system triggers the Data Quality Guard without invoking speculative LLM remediations.

4. Causal Reasoning & Epistemic Audit Design

4-Layer Causal Ontology

The system enforces a strict 4-layer qualitative ontology to prevent conflating symptoms with causes:

4-LAYER CAUSAL ONTOLOGY	
1. ROOT CAUSE	The initiating event (e.g. INTERNAL_RELEASE, EXTERNAL_PROVIDER, INFRASTRUCTURE_FAILURE)
2. AFFECTED SUBSYSTEM	The architectural component experiencing stress (e.g. payment_gateway, checkout_service, inventory_db)
3. PROXIMAL MECHANISM	The physical/technical failure mode (e.g. connection_pool_exhaustion, memory_leak, packet_loss)
4. SYMPTOM KPIs	The observable statistical telemetry anomalies (e.g. gateway_latency_15min, hourly_conversion, revenue)

Example Verified Chain:

INTERNAL_RELEASE (Checkout v4.3) → payment_gateway → connection_pool_exhaustion → gateway_latency_15min (+3.44σ) & conversion drop (-14.2%) .

Stage E6: Deterministic Audit Architecture & Weakest-Link Scoring

Audit scoring is a deterministic function of evidence, rule verdicts, and configured weights. In Stage E6, every candidate hypothesis is audited against **5 deterministic verification rules**:

- timeline** ($w_1 = 0.25$): Validates that initiating event timestamps precede the anomaly window ($t_{evidence} \leq t_{anomaly}$).
- segment_alignment** ($w_2 = 0.20$): Validates that the hypothesis aligns with dimensional cohort concentrations identified in E3.
- kpi_corroboration** ($w_3 = 0.20$): Measures the proportion of anomalous KPIs directly explained by the cited evidence records.
- mechanism_consistency** ($w_4 = 0.20$): Validates ontological compatibility between the proposed mechanism and known domain failure modes.
- contradiction** ($w_5 = 0.15$): Evaluates the presence of direct refuting evidence records in the dossier.

The Weakest-Link Scoring Formula

The engine computes the final score using a weakest-link formula to prevent strong evidence from masking rule failures:

$$\text{rule_score} = \sum_{i=1}^5 w_i \times \text{Multiplier}(\text{Verdict}_i) \quad \text{where Multiplier} \in \{1.0 \text{ (PASS)}, 0.5 \text{ (PARTIAL)}, 0.0 \text{ (FAIL)}\}$$
$$\text{capped_support} = \text{clamp}\left(\frac{\text{support_score}}{2.0}, 0.0, 1.0\right)$$

$$\text{capped_penalty} = \text{clamp}\left(\frac{\text{contradiction_score}}{2.0}, 0.0, 1.0\right)$$

$$\text{final_audit_score} = \text{clamp}\left(\min(\text{capped_support}, \text{rule_score}) - \text{capped_penalty}, 0.0, 1.0\right)$$

The Root-Cause Evidence Gate

Under enterprise audit rules, claiming an INTERNAL_RELEASE or EXTERNAL_PROVIDER root cause requires **direct discriminative evidence** (e.g. git commit logs, CI/CD deployment records, third-party status page logs).

* **If Discriminative Evidence is Missing:** The Root-Cause Gate **FAILS**, and the hypothesis audit verdict is automatically **CAPPED AT MARGINAL (\$0.40\$)**, preventing the system from blaming external vendors or internal software without empirical proof.

AuditVerdict Semantics

- VERIFIED** (≥ 0.70): Deterministically verified causal chain with sufficient evidence and passed Root-Cause Gate.
- MARGINAL** ($0.40 - 0.69$): Partially corroborated hypothesis. Autonomous remediation is blocked; safe diagnostic verification protocols are formulated.
- REJECTED** (< 0.40): Refuted by contradictory evidence or failed verification rules.
- ABSTAIN**: Investigation-level verdict triggered when the top hypothesis is below threshold (< 0.40), the margin between competing hypotheses is too narrow ($\Delta < 0.15$), or data quality/cold-start guards activate.

5. Security, Governance & Role Entitlements

The platform enforces zero-trust data governance across three enterprise persona scopes:

ROLE ENTITLEMENT & VISIBILITY MATRIX			
Persona Scope	Authorized Data Sources		Authorized Operational Levers
Lead Analyst (Full System)	ALL (deployments, logs, SRE, orders, inventory, marketing)		Software Rollback, Traffic Reroute, Canary Diagnostic Verification
CFD / Executive (Financial Scope)	orders, marketing, payment gateway aggregates (No SRE)		Budget Realignment, Promotional Hold, Executive Escalation
Regional Manager (Restricted Scope)	inventory, orders for scoped geographic region (No SRE)		Inventory Rebalance, Regional Promotion, Diagnostic Escalation

Security Controls Implemented

- Fail-Closed Governance:** If a persona attempts to formulate an action using an unallowed lever (or outside their authorization scope), execution is blocked pending architecture review.
- Anti-Hallucination Citation Validation:** All cited evidence IDs are validated against the content-hashed, provenance-tracked database records assembled in E4. Ghost citations are rejected with a 100% score penalty.
- Regional Data Boundary Enforcement:** Slices queries strictly to authorized regions (us-east, us-west, eu-west, ap-south, all), preventing cross-regional telemetry leakage.

6. Outcome Simulation (Stage E8)

Stage E8 generates parametric recovery projections under approved operational interventions:

- **Provenance Label:** Explicitly tagged as `[SIMULATED]` across all API responses and UI components to prevent confusing simulated projections with observed real-world telemetry.
- **Parametric Recovery Model:** Modeled using calibrated exponential decay towards baseline target corridors:

$$y(t) = y_{\text{target}} + (y_{\text{anomaly}} - y_{\text{target}}) \cdot e^{-\lambda t}$$

- **Mean Time to Normalcy (MTTN):** Calculated as the duration required for the metric to recover within $\pm 1.0\sigma$ of historical baseline.
- **Non-Remedial Suppression:** When the engine abstains or issues a non-remedial diagnostic action, recovery curves are suppressed, and an explicit causal disclaimer is attached.

7. Institutional Memory & Precedent Lifecycle (Stage E9)

Stage E9 manages long-term institutional knowledge in a ChromaDB vector store:

- **Lifecycle-Filtered Vector Retrieval:** Precedent queries retrieve historical records, with human-validated records receiving a ranking boost, unvalidated records included under standard similarity weighting, and disputed/suppressed records **strictly excluded from candidate sets**.
- **5-State Precedent Lifecycle:**
 1. `UNVALIDATED` : Newly completed investigation pending peer review (eligible for retrieval under baseline weighting).
 2. `VALIDATED` : Approved by senior analyst / SRE (eligible for retrieval with ranking boost).
 3. `PARTIALLY_VALIDATED` : Confirmed mechanism, modified action directive.
 4. `DISPUTED` : Rejected during human review (**strictly excluded from retrieval**).
 5. `SUPPRESSED` : Deprecated due to architectural redesign or outdated infrastructure (**strictly excluded from retrieval**).
- **Closed-Loop Feedback:** Operators submit structured reviews (`POST /feedback`) to update precedent lifecycle states and prevent error propagation.

8. Data Ingestion Architecture & ChromaDB Population

1. Relational Telemetry Store (PostgreSQL)

- Ingests structured time-series KPI intervals, orders, inventory logs, and payment transactions defined in `etl/schema.sql`.
- Populated via `etl/load_synthetic.py` (with fallback to DuckDB/in-memory SQLite for localized standalone testing).

2. Vector Evidence & Precedent Store (ChromaDB)

The vector store manages two distinct categories of collections in ChromaDB (Port 8000):

- **Scenario Evidence Collections (`evidence_{scenario_id}`):**
- Stores unstructured telemetry evidence: customer support tickets (`support_tickets.csv`), deployment logs (`deployment_log.csv`), and engineering release notes (`data/release_notes/*.txt`).
- Embedded using `bge-m3` via Ollama (`http://localhost:11434`) or ChromaDB's default dense embedding function.

- **Population Commands:**

```
```bash
Embed and load core unstructured evidence records (INC_001)
python etl/load_unstructured.py --scenario-id INC_001 --chroma-host localhost --chroma-port 8000
```



# Seed multi-scenario evidence collections (INC\_002, INC\_004, etc.)

---

```
python etl/seed_scenario_evidence.py --chroma-host localhost --chroma-port 8000
...
```

- **Institutional Memory Collections ( `bi_decisions_precedents` ):**
  - Stores completed incident investigation cases, verified winning hypotheses, and structured recommendations.
  - Automatically populated at runtime by **Engine E9 ( `engines/memory.py` )** after each successful investigation run.
  - Maintains the 5-state validation lifecycle ( `UNVALIDATED` , `VALIDATED` , `PARTIALLY_VALIDATED` , `DISPUTED` , `SUPPRESSED` ).

## 3. Live Streaming Ingestion Boundary

- The current system operates against historical interval batches and staged incident datasets. High-frequency Kafka / event-streaming ingestion pipelines are not implemented in this prototype.
-

## 9. Scenarios Catalog

### Core Validated Reference Scenarios

Scenario ID	Incident Name	Business Situation	Expected Engine Behavior	Final Decision State
INC_001	Payment Gateway Latency Regression	Checkout Service v4.3 release exhausted connection pools, inducing latency spikes (+3.44σ) and conversion drops.	Full deterministic causal chain verified against git release logs and gateway error metrics.	<b>Overall Verdict:</b> <b>VERIFIED</b> Winner: H1 (Score: 71%) Action: Immediate software rollback to v4.2 authorized.
INC_002	Simultaneous Conflicting Causes	Concurrent payment gateway timeouts and aggressive competitor 30% discount.	Narrow margin between competing hypotheses triggers Multi-Causal Conflict Guard.	<b>Overall Verdict:</b> <b>ABSTAIN</b> Winner: None ( H1 = Marginal 80%, H2 = Marginal 80%) Action: Governed diagnostic verification / no operational remediation.
INC_003	Sparse Baseline History	Monitored growth domain has fewer than 14 days of baseline telemetry.	Statistical Cold-Start Guard triggers early bypass (0\$ LLM calls) to prevent false-positive anomaly spikes.   **Overall Verdict: `ABSTAIN`** Winner: `None` Reason: Statistical Cold-Start Guard (<14 intervals).     **`INC_004`**   **ETL Ingestion Pipeline Delay**   Upstream Kafka/Airflow ETL lag causes missing interval batches, mimicking a revenue collapse.   Data Quality Guard detects 0% gateway error rates and flags missing batch artifacts vs real sales loss.   **Overall Verdict: `ABSTAIN`** Winner: `None` Reason: Data Quality Guard (ETL Ingestion Lag).     **`INC_005`**   **Seasonal Demand Pattern**   Standard	z

Scenario ID	Incident Name	Business Situation	Expected Engine Behavior	Final Decision State
			seasonal traffic fluctuations within historical corridor bounds (	
* <b>Token Accounting:</b> Tracks precise prompt tokens ( <code>llm_tokens_in</code> ), completion tokens ( <code>llm_tokens_out</code> ), and total token consumption.				
* <b>Cost Estimator</b> ( <code>llm/cost_estimator.py</code> ): Dynamically computes dollar cost per investigation based on active model rate cards (e.g. Groq Qwen 0.05 / 1M tokens vs. OpenAI GPT-4o). * <b>Deterministic Efficiency:</b> By leveraging deterministic early guards (bypassing LLMs for cold-start and nominal states), the engine eliminates unnecessary LLM inference costs for un-actionable telemetry runs. --- ## 16. Known Limitations * <b>Batch vs Streaming Ingestion:</b> Telemetry is analyzed across fixed historical windows (15\text{m}\$ to 24h); sub-second streaming event ingestion is not implemented in this prototype.				
* <b>Simulated Recovery Projections:</b> Stage E8 recovery trajectories are mathematical exponential approximations and should not be interpreted as guaranteed real-world SLA outcomes.				
* <b>Vector Semantic Scope:</b> Unstructured evidence retrieval depends on the quality and density of ingested log summaries and incident postmortems.				

## 17. Repository Structure

```
BusinessIntelligence.ai/
├─ api/
│ │ └─ __init__.py
│ │ └─ main.py # FastAPI application endpoints (/investigate, /feedback, /health)
├─ config/
│ │ └─ business_denylist.yaml # Prohibited speculative actions & unallowed levers
│ │ └─ domain_semantics.yaml # Domain ontology mappings & mechanism taxonomy
│ │ └─ entitlements.yaml # Role-based persona access policies & source scopes
│ │ └─ evidence_mappings.yaml # Source-to-mechanism correlation weights
│ │ └─ kpi_contracts.yaml # KPI definitions, directions, and baseline corridors
│ │ └─ loader.py # Strict YAML schema validation & configuration loader
│ │ └─ memory_retention.yaml # Vector retention thresholds & invalidation rules
│ │ └─ registry.py # Source and entity registry loader
│ │ └─ scenarios.yaml # Demonstration incident scenario catalog (INC_001 - INC_008)
├─ engines/
│ │ └─ challenge.py # Engine E6: Deterministic Constraint Auditing & Root-Cause Gate
│ │ └─ decision.py # Engine E7: Governed Decision Formulation & Action Directives
│ │ └─ diagnostic.py # Engine E3: Multi-Dimensional Slicing & Decomposition
│ │ └─ evidence.py # Engine E4: Grounded Evidence Dossier Hybrid Retrieval
│ │ └─ hypothesis.py # Engine E5: 4-Layer Causal Ontology Hypothesis Formulation
│ │ └─ kpi_store.py # Engine E1: Time-Series KPI Telemetry Ingestion & Baselines
│ │ └─ memory.py # Engine E9: Vector Institutional Memory & 5-State Lifecycle
│ │ └─ outcome.py # Engine E8: Parametric Exponential Outcome Simulation
│ │ └─ signal.py # Engine E2: Anomaly Detection & Revenue Materiality Ranking
├─ etl/
│ │ └─ generate_scenarios.py # Synthetic incident generator & time-series seed data
│ │ └─ load_synthetic.py # Relational telemetry database loader
│ │ └─ load_unstructured.py # Vector database embedding loader (ChromaDB)
│ │ └─ schema.sql # PostgreSQL relational schema definition
├─ llm/
│ │ └─ cost_estimator.py # Per-run token accounting & LLM inference cost estimator
│ │ └─ provider.py # Groq & Ollama LLM provider client abstractions
│ │ └─ telemetry_wrapper.py # Latency & token monitoring wrapper
├─ pipeline/
│ │ └─ investigate.py # 9-engine execution orchestrator & dependency injector
│ │ └─ telemetry.py # Pipeline telemetry aggregation models
├─ security/
│ │ └─ entitlements.py # Server-side role entitlement & source masking engine
├─ tests/
│ │ └─ test_causal_e6_isolation.py # E6 deterministic scoring isolation test suite
│ │ └─ test_challenge_smoke.py # E6 5-rule constraint verification tests
│ │ └─ test_feedback.py # E9 analyst feedback & lifecycle state transition tests
│ │ └─ test_security.py # Role-based entitlement & access control tests
│ │ └─ test_signal.py # E1/E2 statistical corridor & anomaly detection tests
├─ web/
│ │ └─ src/
│ │ │ └─ components/
│ │ │ │ └─ decision/ # E7 Decision Hero & Abstention Card views
│ │ │ │ └─ engines/ # Workspaces for Stages E1 through E9
│ │ │ │ └─ hypothesis/ # E5/E6 Hypothesis Cards & Rule Scorecard components
│ │ │ │ └─ investigation/ # Pipeline Rail, Feedback Bar, Scenario Selector
│ │ │ │ └─ kpi/ # KPI Signal Grid & telemetry cards
│ │ │ │ └─ layout/ # TopBar, LeftObservePanel, RightBeliefPanel
│ │ │ └─ lib/
│ │ │ │ └─ api.ts # Backend API client
│ │ │ │ └─ narrativeHelpers.ts # 100% dynamic payload narrative presenter
│ │ │ │ └─ utils.ts # Metric formatters & LLM tag cleaners
│ │ │ └─ types/
│ │ │ │ └─ investigation.ts # Complete TypeScript schema definitions
│ │ │ └─ App.tsx # Main application view container
```

```
| | └─ index.css # Design system tokens & dark theme styling
| └─ package.json # Frontend package dependencies & scripts
| └─ vite.config.ts # Vite configuration & proxy definitions
└─ docker-compose.yml # PostgreSQL and ChromaDB container definitions
└─ models.py # Core Python dataclass models and type definitions
└─ requirements.txt # Python backend package dependencies
└─ README.md # Complete technical framework documentation
```

---

## 18. Submission & Technical Review

---

This repository represents the completed Round 2 implementation of **BusinessIntelligence.ai**. It demonstrates:

- \* **Mathematical Causal Grounding:** Confidence scores derived exclusively from deterministic rules rather than stochastic LLM tokens.
- \* **Governed Enterprise Safety:** Automatic abstentions during cold-start, multi-causal ambiguity, and data quality lag.
- \* **Role-Based Security:** Server-side entitlement enforcement across Analyst, CFO, and Regional Manager roles.
- \* **Enterprise-Oriented Analytical Console:** Dark-mode console with sub-second responsive interaction states and verified test suites.

*Submitted for evaluation — BusinessIntelligence.ai Technical Architecture Team.*